Model Development Phase Template

Date	7 July 2024
Team ID	SWTID1720076593
Project Title	Visual Diagnostics: Detecting Tomato Plant Diseases through Leaf Image Analysis
Maximum Marks	10 Marks

Initial Model Training Code, Model Validation and Evaluation Report

The initial model training code will be showcased in the future through a screenshot. The model validation and evaluation report will include a summary and training and validation performance metrics for multiple models, presented through respective screenshots.

Initial Model Training Code (5 marks):

[2]:
```
#import libraries
import tensorflow as tf
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
from tensorflow import keras
```

Training Image preprocessing

[6]:
```
training_set = tf.keras.utils.image_dataset_from_directory(
    'train',
    labels="inferred",
    label_mode="categorical",
    class_names=None,
    color_mode="rgb",
    batch_size=32,
    image_size=(128, 128),
    shuffle=True,
    seed=None,
    validation_split=None,
    subset=None,
    interpolation="bilinear",
    follow_links=False,
    crop_to_aspect_ratio=False
)
```

```
Found 18345 files belonging to 10 classes.
```

Validation Image Preprocessing

[9]:
```
validation_set = tf.keras.utils.image_dataset_from_directory(
    'valid',
    labels="inferred",
    label_mode="categorical",
    class_names=None,
    color_mode="rgb",
    batch_size=32,
    image_size=(128, 128),
    shuffle=True,
    seed=None,
    validation_split=None,
    subset=None,
    interpolation="bilinear",
    follow_links=False,
    crop_to_aspect_ratio=False
)
```

```
Found 4585 files belonging to 10 classes.
```

[11]:
```
training_set
```

[13]:
```
for x,y in training_set:
    print(x,x.shape)
    print(y,y.shape)
    break
```

Building Model

[22]:
```
from keras.layers import Dense,Conv2D,MaxPooling2D,Input,Flatten,Dropout
from keras.models import Sequential
```

[24]:
```
model = Sequential()
```

[26]:
```
model.add(Conv2D(filters=32, kernel_size=3, padding='same', activation='relu', input_shape=[128, 128, 3]))
model.add(Conv2D(filters=32, kernel_size=3, activation='relu'))
model.add(MaxPooling2D(pool_size=2, strides=2))
```

[28]:
```
model.add(Conv2D(filters=64, kernel_size=3, padding='same', activation='relu'))
model.add(Conv2D(filters=64, kernel_size=3, activation='relu'))
model.add(MaxPooling2D(pool_size=2, strides=2))
```

[30]:
```
model.add(Conv2D(filters=128, kernel_size=3, padding='same', activation='relu'))
model.add(Conv2D(filters=128, kernel_size=3, activation='relu'))
model.add(MaxPooling2D(pool_size=2, strides=2))
```

[32]:
```
model.add(Conv2D(filters=256,kernel_size=3,padding='same',activation='relu'))
model.add(Conv2D(filters=256,kernel_size=3,activation='relu'))
model.add(MaxPooling2D(pool_size=2,strides=2))
```

[34]:
```
model.add(Conv2D(filters=512,kernel_size=3,padding='same',activation='relu'))
model.add(Conv2D(filters=512,kernel_size=3,activation='relu'))
model.add(MaxPooling2D(pool_size=2,strides=2))
```

[36]:
```
model.add(Dropout(0.25))
```

[38]:
```
model.add(Flatten())
```

[40]:
```
model.add(Dense(units=1500,activation='relu'))
```

[42]:
```
model.add(Dropout(0.4))
```

[44]:
```
#Output Layer
model.add(Dense(units=10,activation='softmax'))
```

Compiling and Training Phase

[47]:
```
!pip install tf_keras
import os
os.environ['TF_USE_LEGACY_KERAS'] = 'True'
```

[49]:
```
model.compile(optimizer=tf.keras.optimizers.Adam(
    learning_rate=0.0001),loss='categorical_crossentropy',metrics=['accuracy'])
```

[51]:
```
model.summary()
```

Training Model

[54]:
```
training_history = model.fit(x=training_set,validation_data=validation_set,epochs=15)
```

Evaluating Model

[57]:
```
#Training set Accuracy
train_loss, train_acc = model.evaluate(training_set)
print('Training accuracy:', train_acc)
```

```
574/574 ━━━━━━━━━━━━━━━━━━━━ 90s 157ms/step - accuracy: 0.9921 - loss: 0.0233
Training accuracy: 0.9926955699920654
```

[59]:
```
#Validation set Accuracy
val_loss, val_acc = model.evaluate(validation_set)
print('Validation accuracy:', val_acc)
```

```
144/144 ━━━━━━━━━━━━━━━━━━━━ 22s 155ms/step - accuracy: 0.9530 - loss: 0.1473
Validation accuracy: 0.9576880931854248
```

Model Validation and Evaluation Report (5 marks):

Model	Summary	Training and Validation Performance Metrics
Model 1	Model: "sequential_1" Layer (type) — Output Shape — Param # conv2d_2 (Conv2D) — (None, 128, 128, 32) — 896 conv2d_3 (Conv2D) — (None, 126, 126, 32) — 9,248 max_pooling2d_1 (MaxPooling2D) — (None, 63, 63, 32) — 0 conv2d_4 (Conv2D) — (None, 63, 63, 64) — 18,496 conv2d_5 (Conv2D) — (None, 61, 61, 64) — 36,928 max_pooling2d_2 (MaxPooling2D) — (None, 30, 30, 64) — 0 conv2d_6 (Conv2D) — (None, 30, 30, 128) — 73,856 conv2d_7 (Conv2D) — (None, 28, 28, 128) — 147,584 max_pooling2d_3 (MaxPooling2D) — (None, 14, 14, 128) — 0 conv2d_8 (Conv2D) — (None, 14, 14, 256) — 295,168 conv2d_9 (Conv2D) — (None, 12, 12, 256) — 590,080 max_pooling2d_4 (MaxPooling2D) — (None, 6, 6, 256) — 0 conv2d_10 (Conv2D) — (None, 6, 6, 512) — 1,180,160 conv2d_11 (Conv2D) — (None, 4, 4, 512) — 2,359,808 max_pooling2d_5 (MaxPooling2D) — (None, 2, 2, 512) — 0 dropout (Dropout) — (None, 2, 2, 512) — 0 flatten (Flatten) — (None, 2048) — 0 dense (Dense) — (None, 1500) — 3,073,500 dropout_1 (Dropout) — (None, 1500) — 0 dense_1 (Dense) — (None, 10) — 15,010 Total params: 7,800,734 (29.76 MB) Trainable params: 7,800,734 (29.76 MB) Non-trainable params: 0 (0.00 B)	training_history = model.fit(x=training_set,validation_data=validation_set,epochs=15) Epoch 1/15 574/574 — 362s 625ms/step - accuracy: 0.3490 - loss: 1.8109 - val_accuracy: 0.7237 - val_loss: 0.7935 Epoch 2/15 574/574 — 347s 604ms/step - accuracy: 0.7837 - loss: 0.6259 - val_accuracy: 0.8517 - val_loss: 0.4082 Epoch 3/15 574/574 — 360s 628ms/step - accuracy: 0.8638 - loss: 0.4009 - val_accuracy: 0.8866 - val_loss: 0.3325 Epoch 4/15 574/574 — 358s 624ms/step - accuracy: 0.9030 - loss: 0.2816 - val_accuracy: 0.9147 - val_loss: 0.2503 Epoch 5/15 574/574 — 364s 635ms/step - accuracy: 0.9250 - loss: 0.2130 - val_accuracy: 0.9293 - val_loss: 0.2212 Epoch 6/15 574/574 — 348s 607ms/step - accuracy: 0.9377 - loss: 0.1812 - val_accuracy: 0.9514 - val_loss: 0.1474 Epoch 7/15 Epoch 8/15 574/574 — 345s 601ms/step - accuracy: 0.9669 - loss: 0.1001 - val_accuracy: 0.9525 - val_loss: 0.1503 Epoch 9/15 574/574 — 345s 601ms/step - accuracy: 0.9650 - loss: 0.0955 - val_accuracy: 0.9391 - val_loss: 0.1833 Epoch 10/15 574/574 — 346s 603ms/step - accuracy: 0.9709 - loss: 0.0848 - val_accuracy: 0.9535 - val_loss: 0.1352 Epoch 11/15 574/574 — 346s 602ms/step - accuracy: 0.9761 - loss: 0.0725 - val_accuracy: 0.9562 - val_loss: 0.1637 Epoch 12/15 574/574 — 345s 601ms/step - accuracy: 0.9791 - loss: 0.0604 - val_accuracy: 0.9494 - val_loss: 0.1661 Epoch 13/15 574/574 — 345s 601ms/step - accuracy: 0.9749 - loss: 0.0734 - val_accuracy: 0.9418 - val_loss: 0.1942 Epoch 14/15 574/574 — 347s 605ms/step - accuracy: 0.9815 - loss: 0.0541 - val_accuracy: 0.9590 - val_loss: 0.1501 Epoch 15/15 574/574 — 346s 602ms/step - accuracy: 0.9871 - loss: 0.0458 - val_accuracy: 0.9577 - val_loss: 0.1364
Model 2	Model: "sequential_2" Layer (type) — Output Shape — Param # conv2d_8 (Conv2D) — (None, 128, 128, 64) — 1,792 conv2d_9 (Conv2D) — (None, 126, 126, 64) — 36,928 max_pooling2d_4 (MaxPooling2D) — (None, 63, 63, 64) — 0 conv2d_10 (Conv2D) — (None, 63, 63, 128) — 73,856 conv2d_11 (Conv2D) — (None, 61, 61, 128) — 147,584 max_pooling2d_5 (MaxPooling2D) — (None, 30, 30, 128) — 0 conv2d_12 (Conv2D) — (None, 30, 30, 256) — 295,168 conv2d_13 (Conv2D) — (None, 28, 28, 256) — 590,080 max_pooling2d_6 (MaxPooling2D) — (None, 14, 14, 256) — 0 conv2d_14 (Conv2D) — (None, 14, 14, 512) — 1,180,160 conv2d_15 (Conv2D) — (None, 12, 12, 512) — 2,359,808 max_pooling2d_7 (MaxPooling2D) — (None, 6, 6, 512) — 0 dropout_1 (Dropout) — (None, 6, 6, 512) — 0 flatten (Flatten) — (None, 18432) — 0 dense (Dense) — (None, 1500) — 27,649,500 dropout_2 (Dropout) — (None, 1500) — 0 dense_1 (Dense) — (None, 10) — 15,010 Total params: 32,349,886 (123.41 MB) Trainable params: 32,349,886 (123.41 MB) Non-trainable params: 0 (0.00 B)	[61]: training_history = model.fit(x=training_set,validation_data=validation_set,epochs=10) Epoch 1/10 287/287 — 1192s 4s/step - accuracy: 0.1767 - loss: 5.6728 - val_accuracy: 0.4877 - val_loss: 1.4440 Epoch 2/10 287/287 — 1124s 4s/step - accuracy: 0.5387 - loss: 1.3701 - val_accuracy: 0.7252 - val_loss: 0.7970 Epoch 3/10 287/287 — 1124s 4s/step - accuracy: 0.7243 - loss: 0.8081 - val_accuracy: 0.8214 - val_loss: 0.5040 Epoch 4/10 287/287 — 1121s 4s/step - accuracy: 0.7935 - loss: 0.5965 - val_accuracy: 0.8423 - val_loss: 0.4319 Epoch 5/10 287/287 — 1118s 4s/step - accuracy: 0.8625 - loss: 0.3823 - val_accuracy: 0.8554 - val_loss: 0.4279 Epoch 6/10 287/287 — 1117s 4s/step - accuracy: 0.8876 - loss: 0.3243 - val_accuracy: 0.8879 - val_loss: 0.3183 Epoch 7/10 287/287 — 1117s 4s/step - accuracy: 0.9155 - loss: 0.2415 - val_accuracy: 0.9115 - val_loss: 0.2621 Epoch 8/10 287/287 — 2217s 8s/step - accuracy: 0.9288 - loss: 0.2015 - val_accuracy: 0.9400 - val_loss: 0.1750 Epoch 9/10 287/287 — 2792s 10s/step - accuracy: 0.9359 - loss: 0.1920 - val_accuracy: 0.9296 - val_loss: 0.2078 Epoch 10/10 287/287 — 1267s 4s/step - accuracy: 0.9359 - loss: 0.1817 - val_accuracy: 0.9241 - val_loss: 0.2297
Model 3	[9]: model.summary() Model: "sequential" Layer (type) — Output Shape — Param # conv2d (Conv2D) — (None, 128, 128, 128) — 3,584 conv2d_1 (Conv2D) — (None, 126, 126, 128) — 147,584 max_pooling2d (MaxPooling2D) — (None, 63, 63, 128) — 0 conv2d_2 (Conv2D) — (None, 63, 63, 128) — 147,584 conv2d_3 (Conv2D) — (None, 61, 61, 128) — 147,584 max_pooling2d_1 (MaxPooling2D) — (None, 30, 30, 128) — 0 conv2d_4 (Conv2D) — (None, 30, 30, 256) — 295,168 conv2d_5 (Conv2D) — (None, 28, 28, 256) — 590,080 max_pooling2d_2 (MaxPooling2D) — (None, 14, 14, 256) — 0	training_history = model.fit(x=training_set,validation_data=validation_set,epochs=5) Epoch 1/5 144/144 — 1698s 12s/step - accuracy: 0.3355 - loss: 2.3881 - val_accuracy: 0.7210 - val_loss: 0.8021 Epoch 2/5 144/144 — 2974s 20s/step - accuracy: 0.7376 - loss: 0.7630 - val_accuracy: 0.8185 - val_loss: 0.5269 Epoch 3/5 144/144 — 1914s 13s/step - accuracy: 0.8237 - loss: 0.5067 - val_accuracy: 0.8916 - val_loss: 0.3249 Epoch 4/5 144/144 — 1284s 9s/step - accuracy: 0.8764 - loss: 0.3586 - val_accuracy: 0.8995 - val_loss: 0.2970 Epoch 5/5 144/144 — 1036s 7s/step - accuracy: 0.8942 - loss: 0.3022 - val_accuracy: 0.9210 - val_loss: 0.2377

	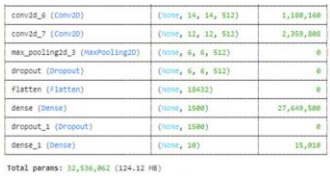 conv2d_6 (Conv2D) (None, 14, 14, 512) 1,180,160 conv2d_7 (Conv2D) (None, 12, 12, 512) 2,359,808 max_pooling2d_3 (MaxPooling2D) (None, 6, 6, 512) 0 dropout (Dropout) (None, 6, 6, 512) 0 flatten (Flatten) (None, 18432) 0 dense (Dense) (None, 1500) 27,649,500 dropout_1 (Dropout) (None, 1500) 0 dense_1 (Dense) (None, 10) 15,010 Total params: 32,536,062 (124.12 MB)	